


Dt:- 16/04/22
9:28 PM

* Regularization

Q: How To Reduce The **Overfitting problem** in Regression?

A:- Ex:- Polynomial Regression \rightarrow Train Accuracy $>$ Test Accuracy



* - Train Data

* - Test Data

Train-error = Low
Test-error = High


* Simple Linear Regression [86%]

* Polynomial Regression [98%]


(Gives best solution for Training) \Rightarrow [Overfitting]
But Less in Test Accuracy

Ans:- "Regularization"

- * Minimizing model complexity
- * Penalizing the loss function
- * Reducing model overfitting [Add more bias to Reduce model variance].



Loss function : $[\sum [y - \hat{y}]^2]_{\min}$
 $[S.S.E]_{\min}$



Variance
Hgh

Bias low

We are adding bias
To Balance [Add "Penalty Term"]

⇒ Three Main types of "Regularization" :

1. L_1 Regularization [Lasso Regression]
2. L_2 Regularization [Ridge Regression]
3. Combining L_1 and L_2 [Elastic Net]

L_2 Regularization
 - Ridge Regression

Ridge regression Advantage is to avoid overfitting problem

Overfitting: performs good in Training data and poor in test Data.


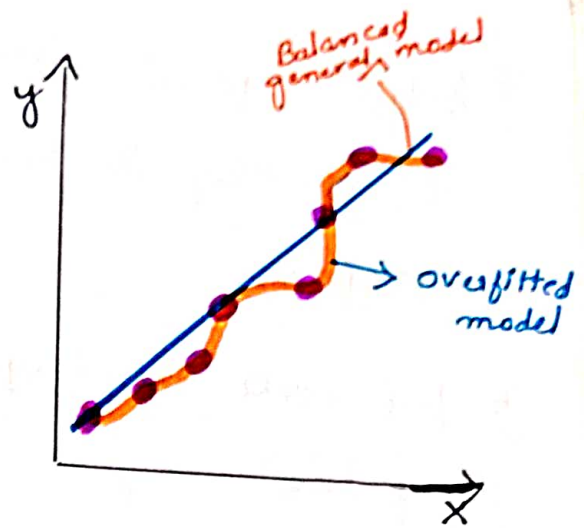
Ridge regression works by applying a penalizing term [Reducing the weights and bias] to overcome overfitting

Least Sum of Squares is applied to obtain best fit line

Since The line passes through the 3 Training data points The Sum of Squared Residuals = 0

But, For Test data, we have More errors. i.e. high

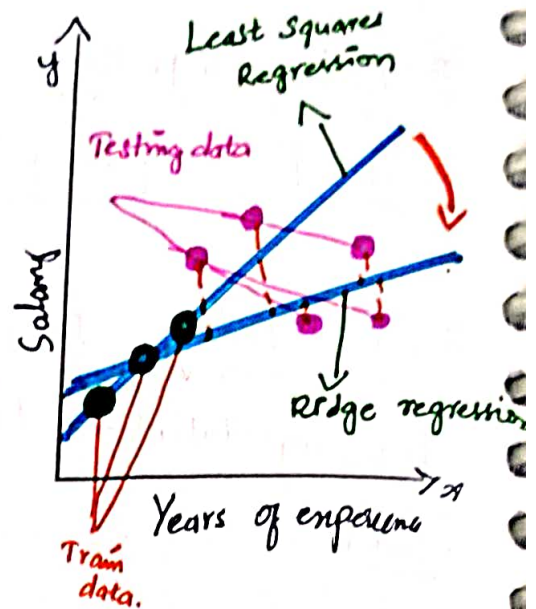
Variance \rightarrow it means that there is a difference in fit (or) (variability) between the training dataset and testing dataset.



Ridge regression works by attempting at increasing the bias to improve variance

This works by changing the slope of line

The model performance might be little poor on Training. But it perform consistently well on both Training & Testing



Slope has been reduced with ridge regression. The model becomes less sensitive to changes in independent variables.

Ex :-

| x | y | $[y - \hat{y}]$ | $[y - \hat{y}]^2$ | $\alpha(\text{slope})^2 = \alpha(\beta_1)^2$ |
|---|---|-----------------|-------------------|--|
| | | | | |

$$\sum (y - \hat{y})^2$$

$$\therefore y = mx + c$$

$$y = \beta_0 + \beta_1 x$$

$$[\sum (y - \hat{y})^2 + \alpha(\text{slope})^2]_{\min}$$

Ex:

$$\begin{bmatrix} S_1 \text{ topper} + S_4 \text{ weak} \\ S_2 \text{ topper} + S_5 \text{ weak} \\ S_3 \text{ topper} + S_6 \text{ weak} \end{bmatrix}_{\min}$$

* Loss function = $[\sum (y - \hat{y})^2]_{\min}$
(For regression line)

* Loss function = $[\sum (y - \hat{y})^2 + \alpha(\text{slope})^2]_{\min}$
(Ridge Regression)
Overall minimum

Least Squares Regression

$\min (\text{Sum of Squared residuals})$
 $[S.S.E]_{\min}$

Ridge regression

$\min (\text{Sum of Squared residuals} + \alpha * \text{slope}^2)$
 $[S.S.E + \alpha (\text{slope}^2)]_{\min}$
 $y = \beta_0 + \beta_1 x$
 not Fixed Value
 Penalty Term

Tuning \Rightarrow which is best one (α) value
 Ex: Radio button.

L_2 regularization adds a penalty equal to the square of the magnitude of coefficient

All coefficients are shrunk by the same factor

does not necessarily eliminate coefficient

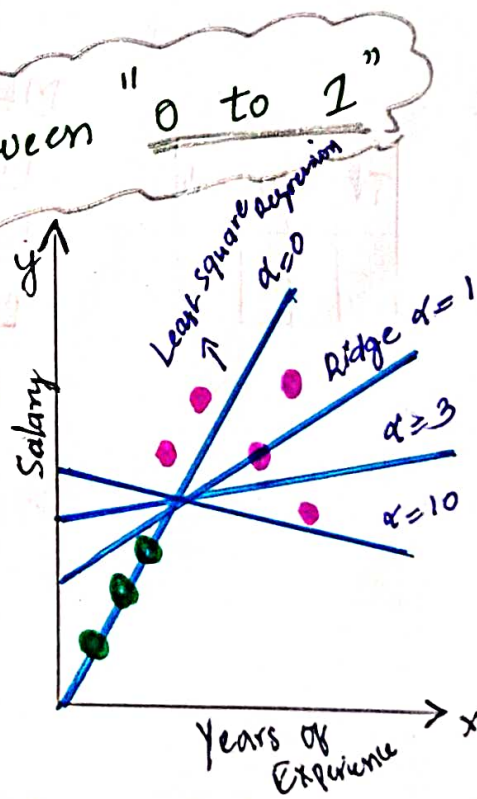
only one all by $\alpha(\beta_1 + \beta_2 + \beta_3)$

Generally

" α " Value should be between "0 to 1"

As Alpha (α) increases, the slope of regression line is reduced and become horizontal.

As Alpha increases, the model becomes less sensitivity to variations of independent variable (# years of Exp)



L2 Regularization

— Lasso Regression.

Lasso regression is similar to Ridge regression.

it works by introducing a bias Term but instead of Squaring Slope The Absolute Value of Slope is added as penalty term

→ Lasso Regression

Min (Sum of Squared residuals + $\alpha * |slope|$)

Apply "modulo" To slope.

Penalty Term

Ex:-

| | | | |
|----|---|-----|------|
| TV | R | Num | Sale |
| | | | |


MLR

PR

Lasso Reg

Ridge reg

Elastic reg



Exam

↓ St

Signature
16/4/22

11:14pm

At: 17/4/22
6:50pm

* Least Square Regression.

Min (sum of The squared Residuals)

* Ridge Regression ✓

Min [sum of the squared residuals + $\alpha * \text{slope}^2$]

* Lasso Regression ✓

Min [sum of squared Residuals + $\alpha * |\text{slope}|$]

Ridge Regression can reduce the slope close to zero (but not exactly zero) but Lasso Regression can reduce the slope to be exactly equal to zero.

HYPER TUNING ∴ [RIDGE, LASSO]

STEP 1

Problem Understanding

STEP 2

Data collection.

```
# df = pd.read_csv("Advertising.csv")
```

```
# df.head()
```

Step 2.2

DataSet Understanding

```
# df.info()
```

Step 2.3

Data Understanding

```
# df.shape
```

Step 3.1

Exploratory Data Analysis [EDA]

```
# df.describe()
```

```
# sns.pairplot(df)
```

```
# df.corr()
```

Data cleaning

Step 3.2

```
# df.isnull().sum()
```

Step 3.3

Data Wrangling

Step 3.4

Train-test-split

```
# X = df[['TV', 'radio', 'newspaper']]  
# Y = df['sales']
```

Always "X"
be
Two dimensional
Array.

```
from sklearn.model_selection import train_test_split  
# X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size = 0.3, Random_state = 29)
```


STEP: 4 MODELLING with default parameters.

→ Ridge Regression

from sklearn.linear_model import Ridge

ridge_model = Ridge()

Help:

Default

Alpha = 1.0

ridge_model.fit(X_train, y_train)

Out: Ridge()

Prediction & Evaluation

test_predictions = ridge_model.predict(X_test)

train_predictions = ridge_model.predict(X_train)

from sklearn.metrics import mean_squared_error

test_rmse = np.sqrt(mean_squared_error(y_test, test_predictions))

train_rmse = np.sqrt(mean_squared_error(y_train, train_predictions))

print("train RMSE :", train_rmse)

print("test RMSE :", test_rmse)

Out:
train RMSE : 1.6408
test RMSE : 1.780

R^2

```
# Print ("Train_R2":, ridge_model.score (x_train, y_train))  
# Print ("Test_R2":, ridge_model.score (x_test, y_test))
```

Out: Train_R² : 0.888
Test_R² : 0.905

⇒ Apply "For Loop" for different Alpha values

for i in range (1,10):

```
# ridge_model = Ridge (alpha = i)  
# ridge_model.fit (x_train, y_train)
```

```
# train_predictions = ridge_model.predict (x_train)  
# test_predictions = ridge_model.predict (x_test)
```

```
# print (i)
```

```
# print ("train-R2:" ridge_model.score (x_train, y_train))  
# print ("test-R2:" ridge_model.score (x_test, y_test))
```

Out:

1. Train-R² : 0.888
Test-R² : 0.905

2. Train-R² : 0.888
Test-R² : 0.905

3. Train-R² : 0.888
Test-R² : 0.905

4. Train-R² : 0.888
Test-R² : 0.905

5. Train-R² : 0.888
Test-R² : 0.905

6. Train-R² : 0.888
Test-R² : 0.905

7. Train-R² : 0.888
Test-R² : 0.905

8. Train-R² : 0.888
Test-R² : 0.905

9. Train-R² : 0.888
Test-R² : 0.905

Signature

⇒ Hyper Parameter tuning

⇓
model Parameter

Ex:- Machine Learning Algorithm : Ridge ()

Help:

Ridge (
 alpha = 1.0,
 *,
 fit_intercept = True,
 normalize = False,
 ...
)

∴ α (Alpha) = hyperparameter.

⇒ **Grid Search CV** (instead of **FOR loop** & Best one & CV)
Sequential order * which is best (given data)
Identifying the best hyperparameter *
cross validation.

MODEL

estimator = Ridge ()

Parameter grid

param_grid = { "alpha": [0, 0.1, 0.2, 0.3, 0.4, 0.5, 1, 2, 10] }

from sklearn.model_selection import GridSearchCV

model_hp = GridSearchCV (estimator, param_grid, fold=5)

model_hp.fit (x_train, y_train)

Out: GridSearchCV (estimator = Ridge (),
 param_grid = { 'alpha': [0, 0.1, 0.2, 0.3, 0.4, 1, 2, 10] })

model_hp. best_params —

Out: {"alpha": 10}

model_hp. best_score —

Out: 0.8831

→ Lasso Regression

MODELING

from sklearn.linear_model import Lasso

lasso_model = Lasso()

lasso_model.fit(x_train, y_train)

Out: Lasso()

Prediction & Evaluation

test_predictions = lasso_model.predict(x_test)

train_predictions = lasso_model.predict(x_train)

from sklearn.metrics import mean_squared_error

train_rmse = np.sqrt(mean_squared_error(y_test, test_predictions))

test_rmse = np.sqrt(mean_squared_error(y_train, train_predictions))

```
# Print ("train RMSE": train_rmse)
# print (" test RMSE: test_rmse)
```

Out: Train RMSE : 1.642
Test RMSE : 1.768

→ Elastic Net Regression

MODELING

```
from sklearn.linear_model import ElasticNet
```

```
# enr_model = ElasticNet()
```

```
# enr_model.fit(x_train, y_train)
```

Out: Elastic Net()

Predictions & Evaluating

```
# train_predictions = enr_model.predict(x_train)
```

```
# test_predictions = enr_model.predict(x_test)
```

```
from sklearn.metrics import mean_squared_error
```

```
# train_rmse = np.sqrt(mean_squared_error(y_train, train_predictions))
```

```
# test_rmse = np.sqrt(mean_squared_error(y_test, test_predictions))
```

```
# Print ("train RMSE:", train_rmse)
```

```
# Print ("test RMSE:", test_rmse)
```


Out:

Train RMSE : 1.6414

Test RMSE : 1.7739

Dataset (B.P.U)

1. Load (import data)
2. Load data (read)
3. Understand data
4. dataset ← Continuous
Discrete
5. EDA (past)
6. Data cleaning
7. Data wrangling
8. Machine learning
9. Evaluation
10. Save model



| | Default | Hyper Tuning Parameter |
|------------------------------|---|---|
| M.L.R (multiple linear) | $\text{Test } R^2 =$ $\text{Train } R^2 =$ $\text{C.V} =$ | Ex: $\text{Test } R^2 =$ $\text{Train } R^2 = 86$ $\text{C.V} = 90$ |
| P.R (Polynomial) | $\text{Test } R^2 =$ $\text{Train } R^2 =$ $\text{C.V} =$ | $\text{Train } R^2 =$ $\text{degree} = ? \quad \text{Test } R^2 = 88$ $\text{C.V} = 90$ |
| Lasso Regression | $\text{Test } R^2 =$ $\text{Train } R^2 =$ $\text{C.V} =$ | $\alpha = ? \quad \text{Train } R^2 =$ $\text{Test } R^2 = 90$ $\text{C.V} = 88$ |
| Ridge regression | $\text{Test } R^2 =$ $\text{Train } R^2 =$ $\text{C.V} =$ | $\alpha = ? \quad \text{Train } R^2 =$ $\text{Test } R^2 =$ $\text{C.V} =$ |
| Elastic Net regression | $\text{Test } R^2 =$ $\text{Train } R^2 =$ $\text{C.V} =$ | $\alpha = ? \quad \text{Train } R^2 =$ $\text{Test } R^2 =$ $\text{C.V} =$ |

R^2 should be high.
after
(C.V) should be considered

Ex: final = (Lasso / $\alpha = 10$) → model.

Signature
18/4/22
2:04 pm